Reintroducing mothur: 10 years later

Patrick D. Schloss[1][†]

† To whom correspondence should be addressed: pschloss@umich.edu

1 Department of Microbiology and Immunology, University of Michigan, Ann Arbor, MI 48109

Mini-review

Abstract

More than 10 years ago, we published the manuscript describing the mothur software package in *Applied and Environmental Microbiology*. Our goal was to create a comprehensive package that allowed users to analyze amplicon sequence data using the most robust methods available. mothur has helped lead the community through the ongoing sequencing revolution and continues to provide this service to the microbial ecology community. Beyond its success and impact on the field, mothur's development exposed a series of observations that are generally translatable across science. Perhaps the observation that stands out the most is that all science is done in the context of prevailing ideas and available technologies. Although it is easy to criticize choices that were made 10 years ago through a modern lens, if we were to wait for all of the possible limitations to be solved before proceeding, science would stall. Even preceding the development of mothur, it was necessary to address the most important problems and work backwards to other problems that limited access to robust sequence analysis tools. At the same time, we strive to expand mothur's capabilities in a data-driven manner to incorporate new ideas and accommodate changes in data and desires of the research community. It has been edifying to see the benefit that a simple set of tools can bring to so many other researchers.

Looking back on scientific journeys can be instructive to others who are overwhelmed at the prospect of looking forward at their careers (1–5). By no means is my scientific journey over, but since 2002 I have been on a journey that I did not realize I was on. Now that the paper introducing the mothur software package is ten years old and has become the most cited paper published by *Applied and Environmental Microbiology* (7) (Figure 1), it is worth stepping back and using the continued development of mothur as a story that has parallels to many other research stories.

I fondly recall preparing a poster for the 2002 meeting of research groups supported by the NSF-supported Microbial Observatories Program. I wanted to triumphantly show that I had sequenced more than 600 16S rRNA gene sequences from a single 0.5-g sample of Alaskan soil. This was greater sequencing depth than anyone else had achieved for a single sample. As I was preparing the poster, I walked into the office of Jo Handelsman, my postdoctoral research advisor at the University of Wisconsin, and laid out the outline for the poster. She asked if I could add one of those "curvy things", a rarefaction curve, to show where I was in sampling the community. Rarefaction curves and attempts to estimate the taxonomic richness of soil had become popular because of the impactful review by Jennifer Hughes and her colleagues (8). Their seminal paper introduced the field to operational taxonomic units (OTUs), rarefaction curves, and richness estimates. I do not recall whether my poster had a rarefaction curve on it, but Jo's question and that review article primed my career.

Introducing DOTUR and friends. When Jo asked me to generate a rarefaction curve for the poster, the request was not trivial. How would I bin the sequences into OTUs? Hughes and her colleagues did it manually and with fewer than 300 sequences. Although I could possibly do that for my 600 sequences, my goal was to generate 1,000 sequences from the sample and to repeat that sampling effort with other samples. I needed something that could be automated. Furthermore, the software that Hughes used to build rarefaction curves, EstimateS (4), required a series of tedious data formatting steps to perform the analyses we were interested in performing. I had found my first problem. How would I assign sequences to OTUs and use that data to estimate the richness and diversity of a sample? The second problem would involve comparing the abundance of OTUs found in one sample to another sample. The solution to the first problem,

3

DOTUR (Distance-based OTUs and Richness), took us two years to develop (9). DOTUR did two things: given a matrix quantifying the genetic distance between pairs of sequences, it would cluster those sequences into OTUs for any distance threshold to define the OTUs and then it would use the frequency of each OTU to calculate a variety of alpha diversity metrics. The solutions to the second problem would come from our work to develop software including ∫-LIBSHUFF (10), SONS (Shared OTUs and Similarity) (11), and TreeClimber (12). Around the same time, Catherine Lozupone and Rob Knight were developing their UniFrac tools to compare communities with a phylogenetic rather than OTU-based approach (13, 14). With these tools, the field of microbial ecology had a quantitative toolbox for describing and comparing microbial communities. Along the way Jo and I would demonstrate the utility of such tools for answering questions like how many OTUs were there in that sample of Alaskan soil and how many sequences were needed to sample each of those OTUs (15)? Where were we in the global bacterial census (16)? How does the word usage of *Goodnight, Moon* compare to that of *Portrait of a Lady* and more importantly how is this relevant to microbial ecology (17)? Most edifying were the more than 2,400 papers that used DOTUR, SONS, TreeClimber, or ∫-LIBSHUFF to facilitate their own research questions (Web of Science, October 1, 2019). Had we waited to solve all of the problems that plagued 16S rRNA gene sequencing, we would still be waiting.

It is important to remember that we knew there were many problems with 16S rRNA gene sequencing. We knew there were biases from extractions and amplification (18–23). We knew there were chimeras (24–27). We knew that bacteria varied in their *rrn* copy number. Generating a distance matrix was a prerequisite to using my tools. This was not trivial, but by cobbling together other tools it was possible. We would assemble, trim and correct Sanger sequence reads using Chromas or STADEN (28), align the sequences using ClustalW (29) or ARB (30), check for chimeras using partial treeing or Bellerephon (27), and calculate a pairwise distance matrix using DNADIST from the PHYLIP package (31). At the time, we knew that we only had a loose concept of a species based on these distances (32). We hoped that an OTU defined as a group of sequences more than 97% similar to each other would be a biologically meaningful unit regardless of whether it fit our notion of a bacterial species. At the time, I felt that the biggest problems that I could solve were how to cluster the sequences into OTUs and how to

4

75 use those clusterings to test our hypotheses. The only tool available at the time that automated

76 the clustering step was FastGroup, which implemented an approximation of the single linkage

77 algorithm (33). The high cost of sequencing was also an impediment to experimentation and

78 analysis in microbial ecology. It was rare for a study design to have experimental replicates so that

79 one could perform a statistical test to compare treatment groups. For example, in our testing we

80 frequently used a dataset comparing Scottish soils from Alison McCaig and colleagues (34). This

81 dataset consisted of two experimental groups, each replicated three times with 45 sequences

82 per replicate. Although great focus has been placed on the depth of sampling afforded by 454

83 and Illumina sequencing, the true benefit of the modern sequencing platforms is the ability to

84 affordably sequence a large number of technical and biological replicates. In my opinion, this

85 expansion in the number of replicates more than makes up for the potential limitations incurred by

86 their shorter read lenght. In spite of the many technical challenges, we had excuses and heuristics

87 to solve problems that served our needs. It is telling that a recent review of "best practices" in

88 generating and analyzing 16S rRNA gene sequences shows that we still have not solved many of

89 these issues and that we have, of course, identified additional problems (35).

90 As we developed these tools, I found a unique niche in microbiology. My undergraduate and

91 graduate training as a biological engineer prepared me to think about research questions

92 from a systems perspective, to think quantitatively, and to understand the value of using

93 computer programs to help solve problems. As an undergraduate student, I learned the Pascal

94 programming language and promptly forgot much of it. Although it was a good language for

95 teaching programming concepts, it did not catch on outside of the classroom. Later, I learned

96 MATLAB. Because it was an expensive commercial programming environment and never caught

97 on with biologists, I also forgot much of it. Even if I forgot the programming syntax of these

98 languages, what learning these languages taught me was the logic and structure of programming.

99 As a postdoc, I would use this background to learn the Perl programming language to better

100 understand how LIBSHUFF (i.e. LIBrary SHUFFle), a tool for comparing the structure of two

101 communities, worked since it was written in Perl (36). After writing my own version of LIBSHUFF,

102 $\int$-LIBSHUFF, and seeing the speed of the version written in C++ by my collaborator, Bret Larget,

103 I converted my Perl version of DOTUR into C++. At the time, the conversion from Perl to C++

seemed like an academic exercise to learn a new language. My Perl version of DOTUR took a minute or so to process the final collection of 1,000 sequences and the C++ version took seconds. Was that really such a big difference? In hindsight, as we now process datasets with tens of millions of sequences, the decision to learn C++ was critical. The ability to pick up computer languages to solve problems, enabled by my prior training in engineering, was a skill that was virtually unheard of in microbiology. Today, researchers without the ability to program are at a significant disadvantage (37).

Introducing mothur. Shortly after DOTUR was published, I received an email from Mitch Sogin, a scientist at the Marine Biology Laboratory (Woods Hole, MA), who asked whether DOTUR could handle more than a million sequences. Without answering his question, I asked where he found a million sequences. Little did I know that his email would represent another pivot in the development of these tools and my career. His group would be the first to use 454 sequencing technology to generate 16S rRNA gene sequences (38). Although DOTUR could assign millions of sequences to OTUs, it was slow and required a significant amount of RAM. As I left my postdoc to start my independent career across the state from Sogin's lab at the University of Massachusetts in Amherst, my plan was to rewrite DOTUR, SONS, $\int$-LIBSHUFF, and TreeClimber for the new world of massively parallelized sequencing. The new tool would become mothur.

Milling about at a poster session at the 2007 ASM General Meeting in Toronto, I ran into Mitch who asked what my plans were for my new lab. I told him that I wanted to make a tool like ARB, a powerful database tool and phylogenetics package (30), but for microbial ecology analysis. His retort was, "You and what army?" Up to that point, I had written every line of code and been answering many emails from people asking for help. He was right, I would need an army. It would be difficult, but I needed to learn to let go and share the development process with someone else. My "army" ended up being Sarah Westcott who has worked on the mothur project from its inception. Today, mothur is nearly 200,000 lines of code and Sarah has touched or written nearly every line of it. Beyond writing and testing mothur's code base, she has become a conduit for many who are trying to learn the tools of microbial ecology. She patiently answers questions via email and on the package's discussion forum (https://forum.mothur.org). The community and I are lucky that Sarah has stayed with the project for more than a decade. To be honest, such dependency on

6

¹³³ a single person makes the project brittle. In hindsight, it would have been better to have developed

¹³⁴ mothur with more of an "army" or team so that there is overlap in people's understanding of how

¹³⁵ mothur works. Although a distributed team approach might work in a software engineering firm, it

¹³⁶ is not practical in most academic environments where there is limited funding. There are certainly

¹³⁷ projects that make this work, but they are rare.

¹³⁸ ***Competition has been good and healthy.*** mothur has not been developed in a vacuum and it

¹³⁹ does not have a monopoly within the field. As indicated above, each of our decisions were made

¹⁴⁰ in the historical context of the field and with constant pressure from others developing their own

¹⁴¹ tools for analyzing 16S rRNA gene sequence data. Competition has been good for mothur and for

¹⁴² the field.

¹⁴³ From the beginning there have been online tools available at the Ribosomal Database Project

¹⁴⁴ (RDP) (39), greengenes (40), and SILVA (41). These allowed users a straightforward method of

¹⁴⁵ comparing their data to those collected in a database. There are two primary downsides to these

¹⁴⁶ tools. First, researchers running the online tool must pay the computational expenses. When their

¹⁴⁷ hardware becomes outdated because it is expensive to replace or maintain, processing times slow

¹⁴⁸ down. Eventually this limitation would result in the termination of the greengenes website. Second,

¹⁴⁹ these platforms provide a one-size-fits-all analysis. These tools only allow a user to analyze 16S

¹⁵⁰ and in some cases 18S rRNA gene sequences. If a user sequences a different gene, then the

¹⁵¹ tool will not serve them. These observations resulted in two design goals we have had with

¹⁵² mothur: bringing the analysis to a user's computer and separating a tool from a specific database.

¹⁵³ For example, we commonly use a sequence alignment method that was originally developed for

¹⁵⁴ greengenes (42), but use a SILVA-based reference alignment because its superior quality (43, 44).

¹⁵⁵ In addition, we offer the naïve Bayesian classifier developed by the RDP (45) and allow users to

¹⁵⁶ train it to any database they want, including customized databases. In both examples, users can

¹⁵⁷ align or classify non-rRNA gene sequence data. As the bioinformatics tools have matured, both

¹⁵⁸ the RDP and SILVA offer integrated pipelines for analyzing large datasets, albeit in one-size-fits-all

¹⁵⁹ black box implementations.

¹⁶⁰ With the growth in popularity of 16S rRNA gene sequencing there has naturally been an

7

expansion in the number of people developing tools to analyze these data. Months after the paper describing mothur was published, the paper describing QIIME was published (46). Over the past 10 years, many have attempted to create analogies comparing the two programs: Pepsi vs Coke, Apple vs Windows, etc. It is never clear which software is which brand and whether the comparisons are meant as a complement or an insult. Regardless, both programs are very popular. From my perspective, most of the differences are cosmetic (http://blog.mothur.org/2016/01/12/mothur-and-qiime/). To me the most meaningful difference between mothur and QIIME is the choice of algorithms used to cluster sequences into OTUs. QIIME's advocacy for open and closed-reference clustering and USEARCH or VSEARCH-based *de novo* clustering results in lower quality OTU assignments relative to the *de novo* clustering algorithms available within mothur (47, 48). QIIME is set of wrapper scripts that help users to transition data between independent packages. For example, with QIIME (through version 1.9.1), it was even possible to run mothur through QIIME. One can also run the naïve Bayesian classifier through QIIME using the original code developed by the RDP. Structuring QIIME as a set of wrappers caused great frustration for many users because there were numerous software dependencies that had to be installed. The benefits included the ability for users to access to a wider set of tools and for developers to tie their tool into the popular software package. Although the QIIME developers would go on to create virtual machines and use packaging tools to simplify installation, these fixes required sophistication by users who we knew struggled with the basics of navigating a command line. In contrast, when a user runs mothur, they are running mothur. The naïve Bayesian classifier code that is in mothur is a rewritten version of the original code. When we rewrite someone's software we do it with an eye to improving performance, access, and utility for non-16S rRNA gene sequence data. For example, while 454 data was popular, PyroNoise was an effective tool for denoising flowgram data (49). Running the original code required a large Linux computer cluster and knowledge of bash and Perl scripting. When we rewrote the code for mothur, we made it accessible to people using any operating system with a simple command interface (i.e. trim.flows and shhh.flows). Our approach requires significant developer effort, but saves considerable user effort. As this benefit is multiplied across thousands of projects, the savings to users has been considerable.

Beyond the large packages like mothur and QIIME, there has been significant growth in standalone software tools for sequence curation (e.g. PyroNoise (49), PANDAseq (50), DADA2 (51)), chimera checking (e.g. UCHIME (52), ChimeraSlayer (53), Perseus (54)), and clustering (e.g. USEARCH (55), VSEARCH (56), Swarm (57)). Where possible and when warranted, we have implemented many of these algorithms directly into mothur. We have also used this diversity of methods to perform head-to-head comparisons. Most notable is the area of clustering algorithms where there have been a large number of algorithms developed without an obvious method to objectively compare them (47, 48, 58, 59). We applied an objective metric, the Matthew's Correlation Coefficient (MCC), to evaluate numerous algorithms for clustering sequences into OTUs. By performing this type of analysis, we were able to objectively compare the algorithms, make recommendations to the field, and develop new algorithms that outperformed the existing ones. Beyond evaluating clustering algorithms, we have also evaluated methods of denoising sequence data (60–62), assessed reference alignments (43, 44), considered the importance of incorporating secondary structure information in alignments (63), quantified the variation along the 16S rRNA gene (44), and compared the statistical hypotheses tested by commonly used tools (64). We have embraced the competition and diversity of all methods being used to analyze amplicon data. This competition forces us to identify the strengths and weaknesses of various methods so that we can make recommendations to other researchers.

mothur's core principles. As mothur has evolved with the needs of the community, several core principles have emerged that direct its development. First, mothur is a free, open source software package. This has been critical in shaping the direction of mothur. We were content for mothur to be an improved combination of DOTUR and SONS and leverage existing tools for other steps. Yet, we quickly appreciated the need for providing other steps in a sequence analysis pipeline to make other tools more accessible. This decision was motivated by learning that the code for greengenes's (42) and ARB/SILVA's aligners were not open source or publicly available. Thus, we realized that such an important functionality needed to be opened to the community (43). More recently, the rejection of closed source, commercial tools can be seen by the broader adoption of open source alternatives. This has been the case with the rising popularity of VSEARCH over USEARCH within the microbial ecology community (55, 56). Related to insuring that mothur's

code is open source, our second core principle is that we maintain transparency to our users. Perhaps a user does not need to interrogate every line of code, but they need to understand what is happening. Many programs, including online workflows, encapsulate large elements of a pipeline in a single command. In contrast, mothur forces the user to specify each step of the pipeline. Although the former approach makes an analysis easier for a beginner, it stifles users that need greater control or understanding of the assumptions at each step. This control over the pipeline has made it easier for researchers to customize databases or adapt the pipeline to analyze non-16S rRNA gene sequence data. Furthermore, we have provided ample instructional materials to teach users how to implement robust pipelines and the theory behind each step through the project's website (https://www.mothur.org). Third, as I mentioned above, there has been a plethora of methods proposed for generating amplicon sequence data, and curating, aligning, checking for chimeras, classifying, and clustering the data. I am proud of the data-driven approach we have taken to comparing these methods. A description of a new method has limited value if it is not benchmarked against other methods or control datasets. Through this core principle and mothur's large reach into the community, we have helped to develop standards in the analysis of 16S rRNA gene sequence data. Fourth, a focus on enabling reproducibility has always been central to the functionality of mothur. From the beginning, mothur's logfiles have represented a transcript of the user's command and outputs. When it became clear that researchers were not submitting their sequence data to the Sequence Read Archive (SRA), we worked with the SRA developers to create a mothur command (make.sra) that creates a package for submitting sequence data through a special mothur portal. A more ambitious project had its seed on April 1, 2013 when we announced a new "function" in mothur: write.paper. The new command required that the user provide a 454 sff file and a journal title or impact factor. With this information, mothur would generate a manuscript. This April Fools' Day joke was poking fun at software that provided an analysis black box but also at many users' sentiments that data analysis should be so cut and dry. A few years later, we revisited this concept in the scope of reproducibility. Why not explicitly script an analysis from downloading data from the SRA through the rendering of a manuscript ready for submission? This idea gave rise to the development of the Riffomonas reproducible research tutorial series that enables researchers to write their own version of write.paper (65). Perhaps the most important core principle is that my research group uses mothur to analyze the

data we generate. This has proven critical as it again represents transparency and hopefully provides confidence to mothur's users that we are not making recommendations that we do not follow ourselves.

Challenges of making open source count. Anyone can post code to GitHub with a permissive license and claim to be an open source software developer. Far more challenging is engaging the target community to make contributions to that code. Frankly, we have struggled to expand the number of people that make contributions to the mothur code base. One challenge we face is that if we looked to others to contribute code to mothur, they would need to know C++. Given the paucity of microbiologists that can program in a compiled language like C++, expecting that community to provide contributors who can write code in a syntax that prizes execution efficiency over developer efficiency was not realistic. In contrast, the QIIME development team could be more distributed because their code base was primarily written in Python, which prizes developer efficiency over execution efficiency. QIIME is a series of wrappers that allow users to execute other developers' code, making the use of a scripting language like Python attractive. Their choices resulted in many tradeoffs that have impacted ease of installation, usability, execution speed, and flexibility. If we were offered funding to rewrite mothur, we would likely rewrite it as an R package that leaned heavily on the R language's C++ interface packages. Of course, such choices are always best in hindsight. When we started developing mothur, the ability to interface between scripting languages like R and Python and C++ code was not as well developed as it is today. For example, the modern version of the Rcpp package was first released in 2009 and its popularity was not immediate (66). The development of mothur has been a product of the environment that it was created in. Although these decisions have largely had positive outcomes, there have been tradeoffs that caused us to sacrifice other goals.

Beyond contributing to the mothur code base, we sought out other ways to include the community as developers. The paper describing mothur included 15 co-authors, most of whom responded to a call to provide a wiki page that described how they used an early version of mothur to analyze a data set. Our vision was that authors might use the mothur wiki to document reproducible workflows for papers using mothur but to also provide instructional materials for others seeking to adapt mothur for their uses (https://www.mothur.org/wiki). Unfortunately, once the incentive

11

of co-authorship was removed, researchers stopped contributing their workflows to the wiki. Again, this vision and the lack of the community's adoption of wikis as a mechanism for reporting workflows was a product of the environment. Although wikis were popular in the late 2000's, they lacked the ability to directly execute the commands that researchers reported. Such technology would not be possible until the creation of IPython notebooks (2011) and R markdown (2012). Another problem with the wiki approach was that potential contributors did not see the wiki as a community resource. I frequently received emails from scientists telling me that there was a typo on a specific page when the intention was that they could correct the typos without my input. We have been more successful in soliciting input and contributions from the user community through the mothur discussion forum and GitHub-based issue tracker. As mothur has matured, we have been dependent on the user community to use these resources to tell us what features they would like to see included in mothur and where the documentation is confusing or incomplete (https://forum.mothur.org). Often we can count on people not directly affiliated with mothur to provide instruction and their own experience to other users on the forum. We are constantly trying to recruit our "army" and are happy to take any contributions we can. Whether the contributions are to the code base, discussion forum, or suggestions for new tools, these contributions have been invaluable to the growth and popularity of mothur.

Failed experiments. If we never failed, we would not be trying hard enough. Over the past decade we have tried a number of experiments to improve the usability and utility of mothur. One of our first experiments was to use mothur to generate standard vector graphic (SVG)-formatted files of heatmaps and Venn diagrams depicting the overlap between microbial communities. Such visuals were helpful for exploring data; however, I quickly realized that I would never put a mothur-generated figure into a manuscript I wrote. Such visuals require far too much customization to be publication-quality. Although QIIME has incorporated visualization tools through the Emperor package (67), the challenge of users taking default values has downsides, especially when those defaults do not follow good data visualization principles. For example, ordinations with black backgrounds and 3-D ordinations in a 2-D medium now litter the literature. Instead, we have encouraged users to use R packages to visualize mothur-generated results using the minimalR instructional materials that I have developed (http://www.riffomonas.org/minimalR/).

A second experiment was the creation of a graphical user interface (GUI) for running mothur. Forcing users to interact with mothur through the command line has been a significant hurdle for many (Figure 3). Unfortunately, the development effort required to create and maintain a GUI is significant and there is limited funding for such efforts. The newest version of QIIME (starting with version 2.0.0) has emphasized interaction with the tools through a GUI (68) and the related QIITA project offers a web-based GUI (69). It remains to be seen how this experiment will go. Another downside of using a GUI is that there is a risk that reproducibility will suffer if users do not have a mechanism to document their mouse clicks. A significant downside for web interfaces is the frequent inability to document or return to old versions of software and databases. As was experienced with greengenes, if the website is terminated, reproducing old analyses becomes impossible. In mothur, documentation of commands and parameter values is explicit in users can provide a file with a list of commands and the software returns a logfile with all commands and output recorded. Given the heightened focus on reproducibility in recent years, we have extended significant effort in developing instructional materials teaching users how to organize, document, and execute reproducible pipelines that allow a user to go from raw sequence data to a compiled manuscript (65, 70). A final example of a failed experiment was a collaboration with programmers through Google Summer of Code to develop commands in mothur that ran the random forest and SVM machine learning algorithms. Similar to the challenges of developing attractive visuals, fitting the algorithms' hyperparameters, testing, and deploying the resulting models require a significant amount of customization. Furthermore, machine learning is an active area of research where methods are still being developed and improved. Thankfully, there are numerous R and Python packages that do a better job of developing these models (71, 72). In each of our "failed" experiments, the real problems were straying from what mothur does well and failing to grasp what we really wanted the innovation to do.

The future. I will continue to develop mothur for as long as other researchers find it useful. One challenge of such a plan is maintaining the funding to support its development. The development of mothur was initially enabled by a subcontract from a Sloan Foundation grant to Mitch Sogin to support his VAMPS (Visualization and Analysis of Microbial Population Structures) initiative. We used that seed funding to secure an NSF grant and then a grant from NIH for tool

development as part of their Human Microbiome Project. Since that project expired in 2013, we have not had funding to specifically support mothur's development. I have been fortunate to have start-up and discretionary funds generated from other projects to help support mothur. Although there is funding for new tools, there appears to be little appetite by funders to support existing tools. Emblematic of this was the NIH program, Big Data To Knowledge (BD2K), which solicited proposals through the program announcement "Extended Development, Hardening and Dissemination of Technologies in Biomedical Computing, Informatics and Big Data Science (PA-14-156)". This opportunity appeared perfect, except that the National Institute of Allergy and Infectious Diseases (NIAID), the primary supporter of microbiome research at NIH, did not participate in the announcement. Tools like mothur are clearly successful, but need funding mechanisms to continue to mature and support the needs of the research community.

As with anything in science, methods become passé. When we first developed mothur, T-RFLP and DGGE were still commonly used. Today it would be hard to argue that data from those methods meaningfully advance a study relative to what one could get using 16S rRNA gene sequence data. Looking forward, many want to claim that amplicon sequencing is today's DGGE. They claim that researchers should instead move on to shotgun metagenomic sequencing. It is important to note that the two methods answer fundamentally different questions. 16S rRNA gene sequence data describes the taxonomic composition, while metagenomic sequence data tells a researcher about the functional potential and genetic diversity of a community. Both tools provide important information, but they cannot easily replace each other. Although metagenomic data does provide highly resolved taxonomic information, its practical limit of detection is at least an order of magnitude higher than that of amplicon data. For example, we analyzed 10,000 16S rRNA sequences from each of about 500 subjects (73). We can think of this as representing about 1,000,000 genome equivalents (10,000 16S rRNA genes/subject x 500 subjects / 5 16S rRNA gene sequences/genome). Assuming a genome is 4 Mbp, this would represent a sequencing depth of 4 Tbp. Although such a sequencing effort is technically possible, the cost of such an endeavor would be considerable and unlikely to be pursued by most researchers. We estimate that generating and sequencing the libraries at the University of Michigan sequencing core would cost approximately $150 per library. The parallel 16S rRNA gene sequences data would cost

approximately $8 per library. Furthermore, analyzing such a large dataset with an approach that captures the full genetic diversity of the community would be financially and technically prohibitive. Going forward, sequencing technologies will continue to evolve to capture longer and more high quality data and there will always be a need for characterizing the taxonomic diversity of microbial communities. With this in mind, there will always be a place for tools like mothur that can analyze amplicon sequence data.

Of course this does not mean that such tools will remain static. We see three key areas that we will continue to help the field to move forward. First, just as we adapted through the transitions from Sanger to 454 to MiSeq and PacBio sequencing platforms (60–62), we must learn whether data from Oxford Nanopore and other developing sequencing technologies can be an alternative sequencing approach that generates sequence data that is the same quality as existing approaches; thus far, the approach has significant shortcomings for sequencing 16S rRNA gene sequences (74). As with the earlier platforms, we must better understand its error profile so that sequencing errors can be corrected. We have learned that moving forward requires that we maintain or improve sequence quality. No doubt, datasets and read lengths will improve, but these advances should not be made at the cost of data quality. Second, with these improvements, we will need to continue to improve our algorithms. We have already seen that attempts to use low quality MiSeq and HiSeq data caused computational problems leading to the creation of open and closed reference clustering methods, which attempted to circumvent those problems (75, 76). Unfortunately, comparative analyses showed that these methods fail relative to *de novo* clustering methods (47, 48). More work is needed to improve reference-based clustering methods so that larger datasets can be analyzed without sacrificing the quality of OTU assignments. Finally, there are ongoing controversies that need further exploration. These include the validity and utility of amplicon sequence variants (77), the wisdom of removing low frequency sequences (78), and methods of identifying and removing contaminant 16S rRNA gene sequences (79, 80). With each of these areas of development, the broader community can count on our same data-driven approach to answer these questions. It is common for researchers to comment that they pick a specific method or deviate from a suggestion because they "like how the data look". When pressed for an objective definition of how they know the data look "right", they

go quiet. Through the use of data where we actually know what looks right and objective metrics of quality, we will continue to base recommendations on data rather than a gut feeling.

Conclusion. In the paper announcing mothur, we commented that the relationship between 16S rRNA gene sequencing and analysis is very much like the Red Queen in Lewis Carroll's book, *Through the Looking-Glass*. Although some disagreed with this analogy (81), I still feel it is apt. The sequencing technology and rapacious appetite of researchers continues to race on. At the same time, bioinformatics tools must adapt to facilitate our research. I am confident that mothur will be up to this exciting challenge. Beyond its utility for analyzing amplicon sequence data, mothur's history provides lessons that are helpful for other projects that hope to develop a long historical arc. First, mothur is a product of its time. We have always sought to solve a current need to the best of our ability with the tools we had at the time. There are certainly caveats to any analysis of 16S rRNA gene sequence data, but if we had waited until those caveats were resolved, the field never would have progressed. Similarly, we made design choices that we probably would not have made had we started the project today. Second, as we have developed mothur, we have attempted to do so in a data-driven approach where we compare multiple methods. It has not merely been enough to propose a new method: we must show that it meaningfully advances the field. Third, through our failures and successes we have learned to focus on what mothur is good at and create products separate from mothur when distinct needs arise. For example, we have learned that mothur should not have a graphical interface or data visualization tool. Instead, we provide instructional materials to teach users how to use the command line interface and other computational skills like programming in R for data visualization. Finally, mothur was born out of a need for automating the analysis of large 16S rRNA gene sequence datasets. It has been refreshing to see the computational skills of the microbial ecology field grow over the past two decades. Looking ahead, we must all take stock of the challenges we face in microbial ecology and how our individual skills and interests can address these challenges to turn them into opportunities.

Acknowledgements

16

References

1. **Lenski RE**. 2017. Experimental evolution and the dynamics of adaptation and genome evolution in microbial populations. The ISME Journal **11**:2181–2194. doi:10.1038/ismej.2017.69.

2. **Smith DK**. 2018. From fundamental supramolecular chemistry to self-assembled nanomaterials and medicines and back again how sam inspired SAMul. Chemical Communications **54**:4743–4760. doi:10.1039/c8cc01753k.

3. **Barbour AG**, **Benach JL**. 2019. Discovery of the lyme disease agent. mBio **10**:e02166–19. doi:10.1128/mbio.02166-19.

4. **Colwell RK**, **Elsensohn JE**. 2014. EstimateS turns 20: Statistical estimation of species richness and shared species from samples, with non-parametric extrapolation. Ecography **37**:609–613. doi:10.1111/ecog.00814.

5. **Glöckner FO**, **Yilmaz P**, **Quast C**, **Gerken J**, **Beccati A**, **Ciuprina A**, **Bruns G**, **Yarza P**, **Peplies J**, **Westram R**, **Ludwig W**. 2017. 25 years of serving the community with ribosomal RNA gene reference databases and tools. Journal of Biotechnology **261**:169–176. doi:10.1016/j.jbiotec.2017.06.1198.

6. **Casadevall A**, **Fang FC**. 2015. (A)Historical science. Infection and Immunity **83**:4460–4464. doi:10.1128/iai.00921-15.

7. **Schloss PD**, **Westcott SL**, **Ryabin T**, **Hall JR**, **Hartmann M**, **Hollister EB**, **Lesniewski RA**, **Oakley BB**, **Parks DH**, **Robinson CJ**, **Sahl JW**, **Stres B**, **Thallinger GG**, **Horn DJV**, **Weber CF**. 2009. Introducing mothur: Open-source, platform-independent, community-supported software for describing and comparing microbial communities. Applied and Environmental Microbiology **75**:7537–7541. doi:10.1128/aem.01541-09.

8. **Hughes JB**, **Hellmann JJ**, **Ricketts TH**, **Bohannan BJM**. 2001. Counting the uncountable: Statistical approaches to estimating microbial diversity. Applied and Environmental Microbiology **67**:4399–4406. doi:10.1128/aem.67.10.4399-4406.2001.

450   9. **Schloss PD**, **Handelsman J**. 2005. Introducing DOTUR, a computer program for defining

451   operational taxonomic units and estimating species richness.   Applied and Environmental

452   Microbiology **71**:1501–1506. doi:10.1128/aem.71.3.1501-1506.2005.

453   10. **Schloss PD**, **Larget BR**, **Handelsman J**. 2004. Integration of microbial ecology and statistics:

454   A test to compare gene libraries.   Applied and Environmental Microbiology **70**:5485–5492.

455   doi:10.1128/aem.70.9.5485-5492.2004.

456   11. **Schloss PD**, **Handelsman J**. 2006. Introducing SONS, a tool for operational taxonomic

457   unit-based comparisons of microbial community memberships and structures.   Applied and

458   Environmental Microbiology **72**:6773–6779. doi:10.1128/aem.00474-06.

459   12.   **Schloss PD**, **Handelsman J**. 2006.   Introducing TreeClimber, a test to compare

460   microbial community structures.   Applied and Environmental Microbiology **72**:2379–2384.

461   doi:10.1128/aem.72.4.2379-2384.2006.

462   13. **Lozupone C**, **Knight R**. 2005. UniFrac: A new phylogenetic method for comparing microbial

463   communities. Applied and Environmental Microbiology **71**:8228–8235. doi:10.1128/aem.71.12.8228-8235.2005.

464   14. **Lozupone CA**, **Hamady M**, **Kelley ST**, **Knight R**. 2007. Quantitative and qualitative diversity

465   measures lead to different insights into factors that structure microbial communities. Applied and

466   Environmental Microbiology **73**:1576–1585. doi:10.1128/aem.01996-06.

467   15. **Schloss PD**, **Handelsman J**. 2006. Toward a census of bacteria in soil. PLoS Computational

468   Biology **2**:e92. doi:10.1371/journal.pcbi.0020092.

469   16. **Schloss PD**, **Handelsman J**. 2004. Status of the microbial census. Microbiology and

470   Molecular Biology Reviews **68**:686–691. doi:10.1128/mmbr.68.4.686-691.2004.

471   17. **Schloss PD**, **Handelsman J**. 2007. The last word: Books as a statistical metaphor for

472   microbial communities. Annual Review of Microbiology **61**:23–34.

473   18. **Zhou J**, **Bruns MA**, **Tiedje JM**. 1996. DNA recovery from soils of diverse composition. Applied

474   and Environmental Microbiology **62**:316–322.

475 19. **Suzuki MT**, **Giovannoni SJ**. 1996. Bias caused by template annealing in the amplification of
476 mixtures of 16S rRNA genes by pcr. Applied and Environmental Microbiology **62**:625–630.

477 20. **Chandler DP**, **Fredrickson JK**, **Brockman FJ**. 1997. Effect of pcr template concentration on
478 the composition and distribution of total community 16S rDNA clone libraries. Molecular Ecology
479 **6**:475–482.

480 21. **Polz MF**, **Cavanaugh CM**. 1998. Bias in template-to-product ratios in multitemplate pcr.
481 Applied and Environmental Microbiology **64**:3724–3730.

482 22. **Wagner A**, **Blackstone N**, **Cartwright P**, **Dick M**, **Misof B**, **Snow P**, **Wagner GP**, **Bartels J**,
483 **Murtha M**, **Pendleton J**. 1994. Surveys of gene families using polymerase chain reaction: PCR
484 selection and pcr drift. Systematic Biology **43**:250–261.

485 23. **Hansen MC**, **Tolker-Nielsen T**, **Givskov M**, **Molin S**. 1998. Biased 16S rDNA pcr
486 amplification caused by interference from dna flanking the template region. FEMS Microbiology
487 Ecology **26**:141–149.

488 24. **Qiu X**, **Wu L**, **Huang H**, **McDonel PE**, **Palumbo AV**, **Tiedje JM**, **Zhou J**. 2001. Evaluation
489 of pcr-generated chimeras, mutations, and heteroduplexes with 16S rRNA gene-based cloning.
490 Applied and Environmental Microbiology **67**:880–887.

491 25. **Komatsoulis GA**, **Waterman MS**. 1997. A new computational method for detection of
492 chimeric 16S rRNA artifacts generated by pcr amplification from mixed bacterial populations.
493 Applied and Environmental Microbiology **63**:2338–2346.

494 26. **Wang G**, **Wang Y**. 1997. Frequency of formation of chimeric molecules as a consequence of
495 pcr coamplification of 16S rRNA genes from mixed bacterial genomes. Applied and Environmental
496 Microbiology **63**:4645–4650.

497 27. **Hugenholtz P**, **Huber T**. 2003. Chimeric 16S rDNA sequences of diverse origin are
498 accumulating in the public databases. International Journal of Systematic and Evolutionary
499 Microbiology **53**:289–293.

28. **Bonfield JK**, **Smith KF**, **Staden R**. 1995. A new dna sequence assembly program. Nucleic Acids Research **23**:4992–4999.

29. **Thompson JD**, **Higgins DG**, **Gibson TJ**. 1994. CLUSTAL w: Improving the sensitivity of progressive multiple sequence alignment through sequence weighting, position-specific gap penalties and weight matrix choice. Nucleic Acids Research **22**:4673–4680. doi:10.1093/nar/22.22.4673.

30. **Ludwig W**. 2004. ARB: A software environment for sequence data. Nucleic Acids Research **32**:1363–1371. doi:10.1093/nar/gkh293.

31. **Felsenstein J**. 1989. PHYLIP - phylogeny inference package. Cladistics **5**:164–166.

32. **Stackebrandt E**, **Goebel BM**. 1994. Taxonomic note: A place for DNA-DNA reassociation and 16S rRNA sequence analysis in the present species definition in bacteriology. International Journal of Systematic and Evolutionary Microbiology **44**:846–849. doi:10.1099/00207713-44-4-846.

33. **Seguritan V**, **Rohwer F**. 2001. FastGroup: A program to dereplicate libraries of 16S rDNA sequences. BMC Bioinformatics **2**:9. doi:10.1186/1471-2105-2-9.

34. **McCaig AE**, **Glover LA**, **Prosser JI**. 1999. Molecular analysis of bacterial community structure and diversity in unimproved and improved upland grass pastures. Applied and Environmental Microbiology **65**:1721–1730.

35. **Pollock J**, **Glendinning L**, **Wisedchanwet T**, **Watson M**. 2018. The madness of microbiome: Attempting to find consensus "Best practice" for 16S microbiome studies. Applied and Environmental Microbiology **84**:e02627–17. doi:10.1128/aem.02627-17.

36. **Singleton DR**, **Furlong MA**, **Rathbun SL**, **Whitman WB**. 2001. Quantitative comparisons of 16S rRNA gene sequence libraries from environmental samples. Applied and Environmental Microbiology **67**:4374–4376. doi:10.1128/aem.67.9.4374-4376.2001.

37. **Carey MA**, **Papin JA**. 2018. Ten simple rules for biologists learning to program. PLOS Computational Biology **14**:e1005871. doi:10.1371/journal.pcbi.1005871.

38. **Sogin ML**, **Morrison HG**, **Huber JA**, **Welch DM**, **Huse SM**, **Neal PR**, **Arrieta JM**, **Herndl GJ**.

21

2006. Microbial diversity in the deep sea and the underexplored "rare biosphere". Proceedings of the National Academy of Sciences **103**:12115–12120. doi:10.1073/pnas.0605127103.

39. **Cole JR**, **Wang Q**, **Fish JA**, **Chai B**, **McGarrell DM**, **Sun Y**, **Brown CT**, **Porras-Alfaro A**, **Kuske CR**, **Tiedje JM**. 2013. Ribosomal database project: Data and tools for high throughput rRNA analysis. Nucleic Acids Research **42**:D633–D642. doi:10.1093/nar/gkt1244.

40. **DeSantis TZ**, **Hugenholtz P**, **Larsen N**, **Rojas M**, **Brodie EL**, **Keller K**, **Huber T**, **Dalevi D**, **Hu P**, **Andersen GL**. 2006. Greengenes, a chimera-checked 16S rRNA gene database and workbench compatible with ARB. Applied and Environmental Microbiology **72**:5069–5072. doi:10.1128/aem.03006-05.

41. **Yilmaz P**, **Parfrey LW**, **Yarza P**, **Gerken J**, **Pruesse E**, **Quast C**, **Schweer T**, **Peplies J**, **Ludwig W**, **Glöckner FO**. 2013. The SILVA and "All-species living tree project (LTP)" taxonomic frameworks. Nucleic Acids Research **42**:D643–D648. doi:10.1093/nar/gkt1209.

42. **DeSantis TZ**, **Hugenholtz P**, **Keller K**, **Brodie EL**, **Larsen N**, **Piceno YM**, **Phan R**, **Andersen GL**. 2006. NAST: A multiple sequence alignment server for comparative analysis of 16S rRNA genes. Nucleic Acids Research **34**:W394–W399. doi:10.1093/nar/gkl244.

43. **Schloss PD**. 2009. A high-throughput DNA sequence aligner for microbial ecology studies. PLoS ONE **4**:e8230. doi:10.1371/journal.pone.0008230.

44. **Schloss PD**. 2010. The effects of alignment quality, distance calculation method, sequence filtering, and region on the analysis of 16S rRNA gene-based studies. PLoS Computational Biology **6**:e1000844. doi:10.1371/journal.pcbi.1000844.

45. **Wang Q**, **Garrity GM**, **Tiedje JM**, **Cole JR**. 2007. Naive bayesian classifier for rapid assignment of rRNA sequences into the new bacterial taxonomy. Applied and Environmental Microbiology **73**:5261–5267. doi:10.1128/aem.00062-07.

46. **Caporaso JG**, **Kuczynski J**, **Stombaugh J**, **Bittinger K**, **Bushman FD**, **Costello EK**, **Fierer N**, **Peña AG**, **Goodrich JK**, **Gordon JI**, **Huttley GA**, **Kelley ST**, **Knights D**, **Koenig JE**, **Ley RE**, **Lozupone CA**, **McDonald D**, **Muegge BD**, **Pirrung M**, **Reeder J**, **Sevinsky JR**,

Turnbaugh PJ, **Walters WA**, **Widmann J**, **Yatsunenko T**, **Zaneveld J**, **Knight R**. 2010. QIIME allows analysis of high-throughput community sequencing data. Nature Methods **7**:335–336. doi:10.1038/nmeth.f.303.

47. **Westcott SL**, **Schloss PD**. 2015. De novo clustering methods outperform reference-based methods for assigning 16S rRNA gene sequences to operational taxonomic units. PeerJ **3**:e1487. doi:10.7717/peerj.1487.

48. **Westcott SL**, **Schloss PD**. 2017. OptiClust, an improved method for assigning amplicon-based sequence data to operational taxonomic units. mSphere **2**:e00073–17. doi:10.1128/mspheredirect.00073-17.

49. **Quince C**, **Lanzén A**, **Curtis TP**, **Davenport RJ**, **Hall N**, **Head IM**, **Read LF**, **Sloan WT**. 2009. Accurate determination of microbial diversity from 454 pyrosequencing data. Nature Methods **6**:639–641. doi:10.1038/nmeth.1361.

50. **Masella AP**, **Bartram AK**, **Truszkowski JM**, **Brown DG**, **Neufeld JD**. 2012. PANDAseq: Paired-end assembler for illumina sequences. BMC Bioinformatics **13**:31. doi:10.1186/1471-2105-13-31.

51. **Callahan BJ**, **McMurdie PJ**, **Rosen MJ**, **Han AW**, **Johnson AJA**, **Holmes SP**. 2016. DADA2: High-resolution sample inference from illumina amplicon data. Nature Methods **13**:581–583. doi:10.1038/nmeth.3869.

52. **Edgar RC**, **Haas BJ**, **Clemente JC**, **Quince C**, **Knight R**. 2011. UCHIME improves sensitivity and speed of chimera detection. Bioinformatics **27**:2194–2200. doi:10.1093/bioinformatics/btr381.

53. **Haas BJ**, **Gevers D**, **Earl AM**, **Feldgarden M**, **Ward DV**, **Giannoukos G**, **Ciulla D**, **Tabbaa D**, **Highlander SK**, **Sodergren E**, **Methe B**, **DeSantis TZ**, **Petrosino JF**, **Knight R**, **and BWB**. 2011. Chimeric 16S rRNA sequence formation and detection in Sanger and 454-pyrosequenced PCR amplicons. Genome Research **21**:494–504. doi:10.1101/gr.112730.110.

54. **Quince C**, **Lanzen A**, **Davenport RJ**, **Turnbaugh PJ**. 2011. Removing noise from pyrosequenced amplicons. BMC Bioinformatics **12**:38. doi:10.1186/1471-2105-12-38.

576 55.   **Edgar RC**. 2010.   Search and clustering orders of magnitude faster than BLAST.

577 Bioinformatics **26**:2460–2461. doi:10.1093/bioinformatics/btq461.

578 56.   **Rognes T**, **Flouri T**, **Nichols B**, **Quince C**, **Mahé F**. 2016.   VSEARCH: A versatile open

579 source tool for metagenomics. PeerJ **4**:e2584. doi:10.7717/peerj.2584.

580 57. **Mahé F**, **Rognes T**, **Quince C**, **Vargas C de**, **Dunthorn M**. 2015. Swarm v2: Highly-scalable

581 and high-resolution amplicon clustering. PeerJ **3**:e1420. doi:10.7717/peerj.1420.

582 58.   **Schloss PD**, **Westcott SL**. 2011.   Assessing and improving methods used in operational

583 taxonomic unit-based approaches for 16S rRNA gene sequence analysis.   Applied and

584 Environmental Microbiology **77**:3219–3226. doi:10.1128/aem.02810-10.

585 59.   **Schloss PD**. 2016. Application of a database-independent approach to assess the quality of

586 operational taxonomic unit picking methods. mSystems **1**:e00027–16. doi:10.1128/msystems.00027-16.

587 60.   **Kozich JJ**, **Westcott SL**, **Baxter NT**, **Highlander SK**, **Schloss PD**. 2013.   Development of

588 a dual-index sequencing strategy and curation pipeline for analyzing amplicon sequence data on

589 the MiSeq illumina sequencing platform. Applied and Environmental Microbiology **79**:5112–5120.

590 doi:10.1128/aem.01043-13.

591 61.   **Schloss PD**, **Gevers D**, **Westcott SL**. 2011.   Reducing the effects of PCR amplification and

592 sequencing artifacts on 16S rRNA-based studies. PLoS ONE **6**:e27310. doi:10.1371/journal.pone.0027310.

593 62.   **Schloss PD**, **Jenior ML**, **Koumpouras CC**, **Westcott SL**, **Highlander SK**. 2016. Sequencing

594 16S rRNA gene fragments using the PacBio SMRT DNA sequencing system.   PeerJ **4**:e1869.

595 doi:10.7717/peerj.1869.

596 63.   **Schloss PD**. 2012.   Secondary structure improves OTU assignments of 16S rRNA gene

597 sequences. The ISME Journal **7**:457–460. doi:10.1038/ismej.2012.102.

598 64.   **Schloss PD**. 2008. Evaluating different approaches that test whether microbial communities

599 have the same structure. The ISME Journal **2**:265–275. doi:10.1038/ismej.2008.5.

600 65.   **Schloss PD**. 2018.   The riffomonas reproducible research tutorial series.   Journal of Open

Source Education **1**:13. doi:10.21105/jose.00013.

66. **Eddelbuettel D**, **François R**. 2011. Rcpp: Seamless R and C++ integration. Journal of Statistical Software **40**:1–18. doi:10.18637/jss.v040.i08.

67. **Vázquez-Baeza Y**, **Pirrung M**, **Gonzalez A**, **Knight R**. 2013. EMPeror: A tool for visualizing high-throughput microbial community data. GigaScience **2**:16. doi:10.1186/2047-217x-2-16.

68. **Bolyen E**, **Rideout JR**, **Dillon MR**, **Bokulich NA**, **Abnet CC**, **Al-Ghalith GA**, **Alexander H**, **Alm EJ**, **Arumugam M**, **Asnicar F**, **Bai Y**, **Bisanz JE**, **Bittinger K**, **Brejnrod A**, **Brislawn CJ**, **Brown CT**, **Callahan BJ**, **Caraballo-Rodrguez AM**, **Chase J**, **Cope EK**, **Silva RD**, **Diener C**, **Dorrestein PC**, **Douglas GM**, **Durall DM**, **Duvallet C**, **Edwardson CF**, **Ernst M**, **Estaki M**, **Fouquier J**, **Gauglitz JM**, **Gibbons SM**, **Gibson DL**, **Gonzalez A**, **Gorlick K**, **Guo J**, **Hillmann B**, **Holmes S**, **Holste H**, **Huttenhower C**, **Huttley GA**, **Janssen S**, **Jarmusch AK**, **Jiang L**, **Kaehler BD**, **Kang KB**, **Keefe CR**, **Keim P**, **Kelley ST**, **Knights D**, **Koester I**, **Kosciolek T**, **Kreps J**, **Langille MGI**, **Lee J**, **Ley R**, **Liu Y-X**, **Loftfield E**, **Lozupone C**, **Maher M**, **Marotz C**, **Martin BD**, **McDonald D**, **McIver LJ**, **Melnik AV**, **Metcalf JL**, **Morgan SC**, **Morton JT**, **Naimey AT**, **Navas-Molina JA**, **Nothias LF**, **Orchanian SB**, **Pearson T**, **Peoples SL**, **Petras D**, **Preuss ML**, **Pruesse E**, **Rasmussen LB**, **Rivers A**, **Robeson MS**, **Rosenthal P**, **Segata N**, **Shaffer M**, **Shiffer A**, **Sinha R**, **Song SJ**, **Spear JR**, **Swafford AD**, **Thompson LR**, **Torres PJ**, **Trinh P**, **Tripathi A**, **Turnbaugh PJ**, **Ul-Hasan S**, **Hooft JJJ van der**, **Vargas F**, **Vázquez-Baeza Y**, **Vogtmann E**, **Hippel M von**, **Walters W**, **Wan Y**, **Wang M**, **Warren J**, **Weber KC**, **Williamson CHD**, **Willis AD**, **Xu ZZ**, **Zaneveld JR**, **Zhang Y**, **Zhu Q**, **Knight R**, **Caporaso JG**. 2019. Reproducible, interactive, scalable and extensible microbiome data science using QIIME 2. Nature Biotechnology **37**:852–857. doi:10.1038/s41587-019-0209-9.

69. **Gonzalez A**, **Navas-Molina JA**, **Kosciolek T**, **McDonald D**, **Vázquez-Baeza Y**, **Ackermann G**, **DeReus J**, **Janssen S**, **Swafford AD**, **Orchanian SB**, **Sanders JG**, **Shorenstein J**, **Holste H**, **Petrus S**, **Robbins-Pianka A**, **Brislawn CJ**, **Wang M**, **Rideout JR**, **Bolyen E**, **Dillon M**, **Caporaso JG**, **Dorrestein PC**, **Knight R**. 2018. Qiita: Rapid, web-enabled microbiome meta-analysis. Nature Methods **15**:796–798. doi:10.1038/s41592-018-0141-9.

70. **Schloss PD**. 2018. Identifying and overcoming threats to reproducibility, replicability, robustness, and generalizability in microbiome research. mBio **9**:e00525–18. doi:10.1128/mbio.00525-18.

71. **Paszke A**, **Gross S**, **Chintala S**, **Chanan G**, **Yang E**, **DeVito Z**, **Lin Z**, **Desmaison A**, **Antiga L**, **Lerer A**. 2017. Automatic differentiation in PyTorch. *In* NIPS autodiff workshop.

72. **Kuhn M**. 2008. Building predictive models in R using the caret package. Journal of Statistical Software, Articles **28**:1–26. doi:10.18637/jss.v028.i05.

73. **Baxter NT**, **Ruffin MT**, **Rogers MAM**, **Schloss PD**. 2016. Microbiota-based model improves the sensitivity of fecal immunochemical test for detecting colonic lesions. Genome Medicine **8**:37. doi:10.1186/s13073-016-0290-3.

74. **Calus ST**, **Ijaz UZ**, **Pinto AJ**. 2018. NanoAmpli-seq: A workflow for amplicon sequencing for mixed microbial communities on the nanopore sequencing platform. GigaScience **7**:12. doi:10.1093/gigascience/giy140.

75. **Navas-Molina JA**, **Peralta-Sánchez JM**, **González A**, **McMurdie PJ**, **Vázquez-Baeza Y**, **Xu Z**, **Ursell LK**, **Lauber C**, **Zhou H**, **Song SJ**, **Huntley J**, **Ackermann GL**, **Berg-Lyons D**, **Holmes S**, **Caporaso JG**, **Knight R**. 2013. Advancing our understanding of the human microbiome using QIIME, pp. 371–444. *In* Methods in Enzymology. Elsevier.

76. **Rideout JR**, **He Y**, **Navas-Molina JA**, **Walters WA**, **Ursell LK**, **Gibbons SM**, **Chase J**, **McDonald D**, **Gonzalez A**, **Robbins-Pianka A**, **Clemente JC**, **Gilbert JA**, **Huse SM**, **Zhou H-W**, **Knight R**, **Caporaso JG**. 2014. Subsampled open-reference clustering creates consistent, comprehensive OTU definitions and scales to billions of sequences. PeerJ **2**:e545. doi:10.7717/peerj.545.

77. **Callahan BJ**, **McMurdie PJ**, **Holmes SP**. 2017. Exact sequence variants should replace operational taxonomic units in marker-gene data analysis. The ISME Journal **11**:2639–2643. doi:10.1038/ismej.2017.119.

78. **Bokulich NA**, **Subramanian S**, **Faith JJ**, **Gevers D**, **Gordon JI**, **Knight R**, **Mills DA**, **Caporaso JG**. 2012. Quality-filtering vastly improves diversity estimates from illumina amplicon

sequencing. Nature Methods **10**:57–59. doi:10.1038/nmeth.2276.

79. **Salter SJ**, **Cox MJ**, **Turek EM**, **Calus ST**, **Cookson WO**, **Moffatt MF**, **Turner P**, **Parkhill J**, **Loman NJ**, **Walker AW**. 2014. Reagent and laboratory contamination can critically impact sequence-based microbiome analyses. BMC Biology **12**:87. doi:10.1186/s12915-014-0087-z.

80. **Davis NM**, **Proctor DM**, **Holmes SP**, **Relman DA**, **Callahan BJ**. 2018. Simple statistical identification and removal of contaminant sequences in marker-gene and metagenomics data. Microbiome **6**:1. doi:10.1186/s40168-018-0605-2.

81. **Caporaso JG**, **Lauber CL**, **Walters WA**, **Berg-Lyons D**, **Lozupone CA**, **Turnbaugh PJ**, **Fierer N**, **Knight R**. 2010. Global patterns of 16S rRNA diversity at a depth of millions of sequences per sample. Proceedings of the National Academy of Sciences **108**:4516–4522. doi:10.1073/pnas.1000080107.

Figure 1. mothur has consistently been a popular software package over the past ten years with more than 8,800 citations. Citation data taken from the Web of Science (https: //www.webofscience.com) on October 1, 2019. The gray line segment depicts the projected number of citations for 2019 based on the current number of citations for the year and historical trends.

28

Figure 2. The mothur homepage. From the mothur home page at www.mothur.org, users can download mothur, access a user forum, navigate a wiki with extensive documentation, find blog posts that provide additional examples of how to use mothur, join the mothur facebook group, and subscribe to the mothur mailing list.

```
Mac version

Using ReadLine
mothur v.1.44.0
Last updated: 10/02/2019
by
Patrick D. Schloss

Department of Microbiology & Immunology

University of Michigan
http://www.mothur.org

When using, please cite:
Schloss, P.D., et al., Introducing mothur: Open-source, platform-independent, comm
unity-supported software for describing and comparing microbial communities. Appl
Environ Microbiol, 2009. 75(23):7537-41.

Distributed under the GNU General Public License

Type 'help()' for information on the commands that are available

For questions and analysis support, please visit our forum at https://forum.mothur
.org

Type 'quit()' to exit program

[NOTE]: Setting random seed to 19760620.

Interactive Mode


mothur > ▌
```

Figure 3. The start up window when running mothur in Mac OS X in the interactive mode. mothur can also be run on Windows or Linux. In the interactive mode users enter individual commands at the mothur prompt. Alternatively, users may run mothur by supplying commands from the command line or using batch scripts.